

Hardware and Interrupts

Introduction to GPIO



Santosh Sam Koshy
Principal Technical Officer
C-DAC Hyderabad

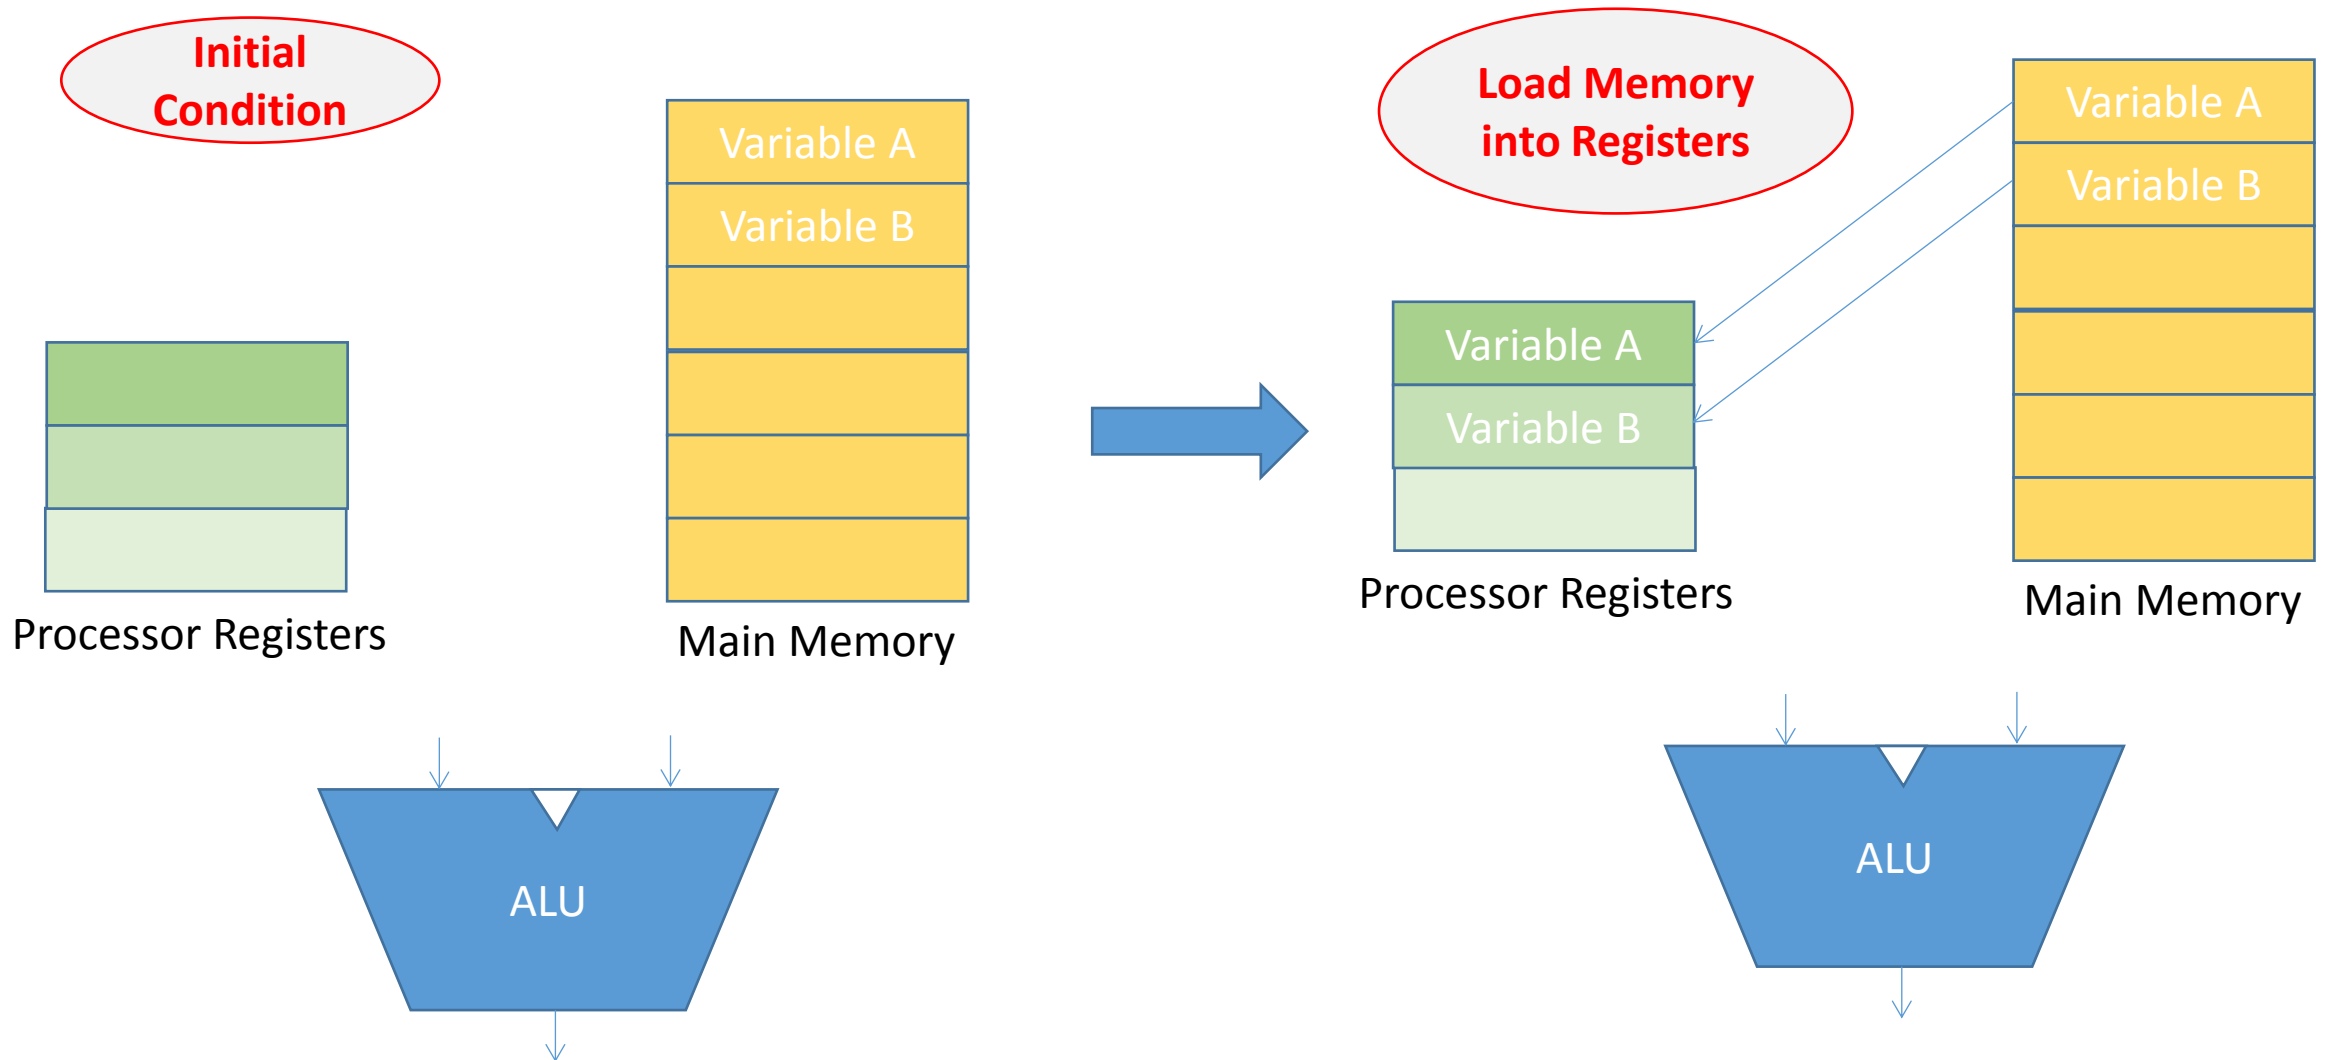
Content

- Device Description
- Memory mapped IO and IO mapped IO
- Handling Interrupts in the kernel
- BBB Header Information
- GPIO Functions
- GPIO as an Input Device
- GPIO as an Output Device

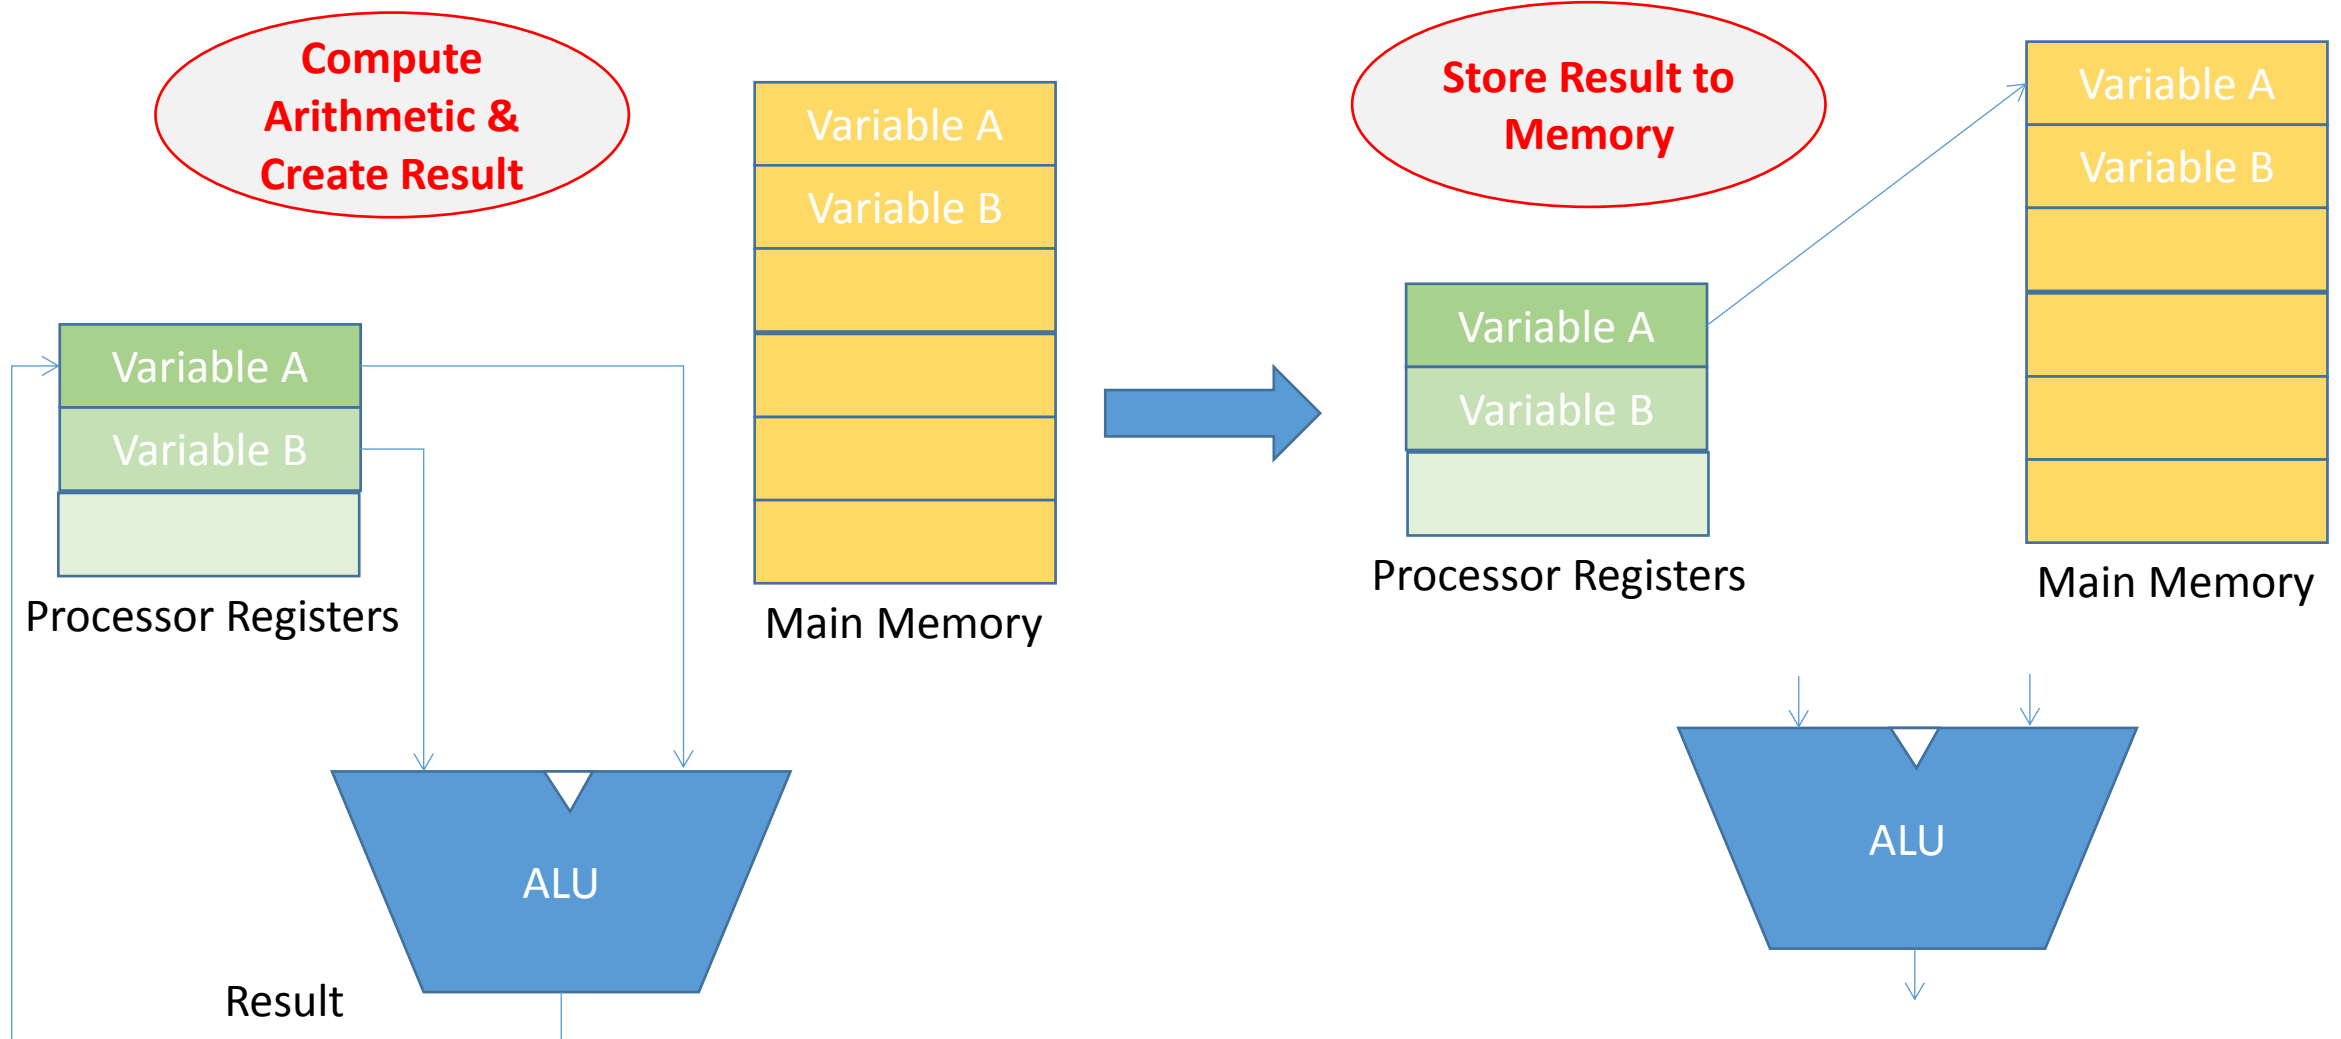
Device Memory

- Peripheral devices are controlled by writing and reading its registers
- Registers
 - Data Registers (Reading and Writing)
 - Control Registers (Configuring the device. Supports R&W Operations)
 - Status Registers (Maintains state of the device. Generally Read Only)
- Every peripheral has a number of such registers, generally mapped contiguously in the Processor's Memory Address Space or IO Address Space

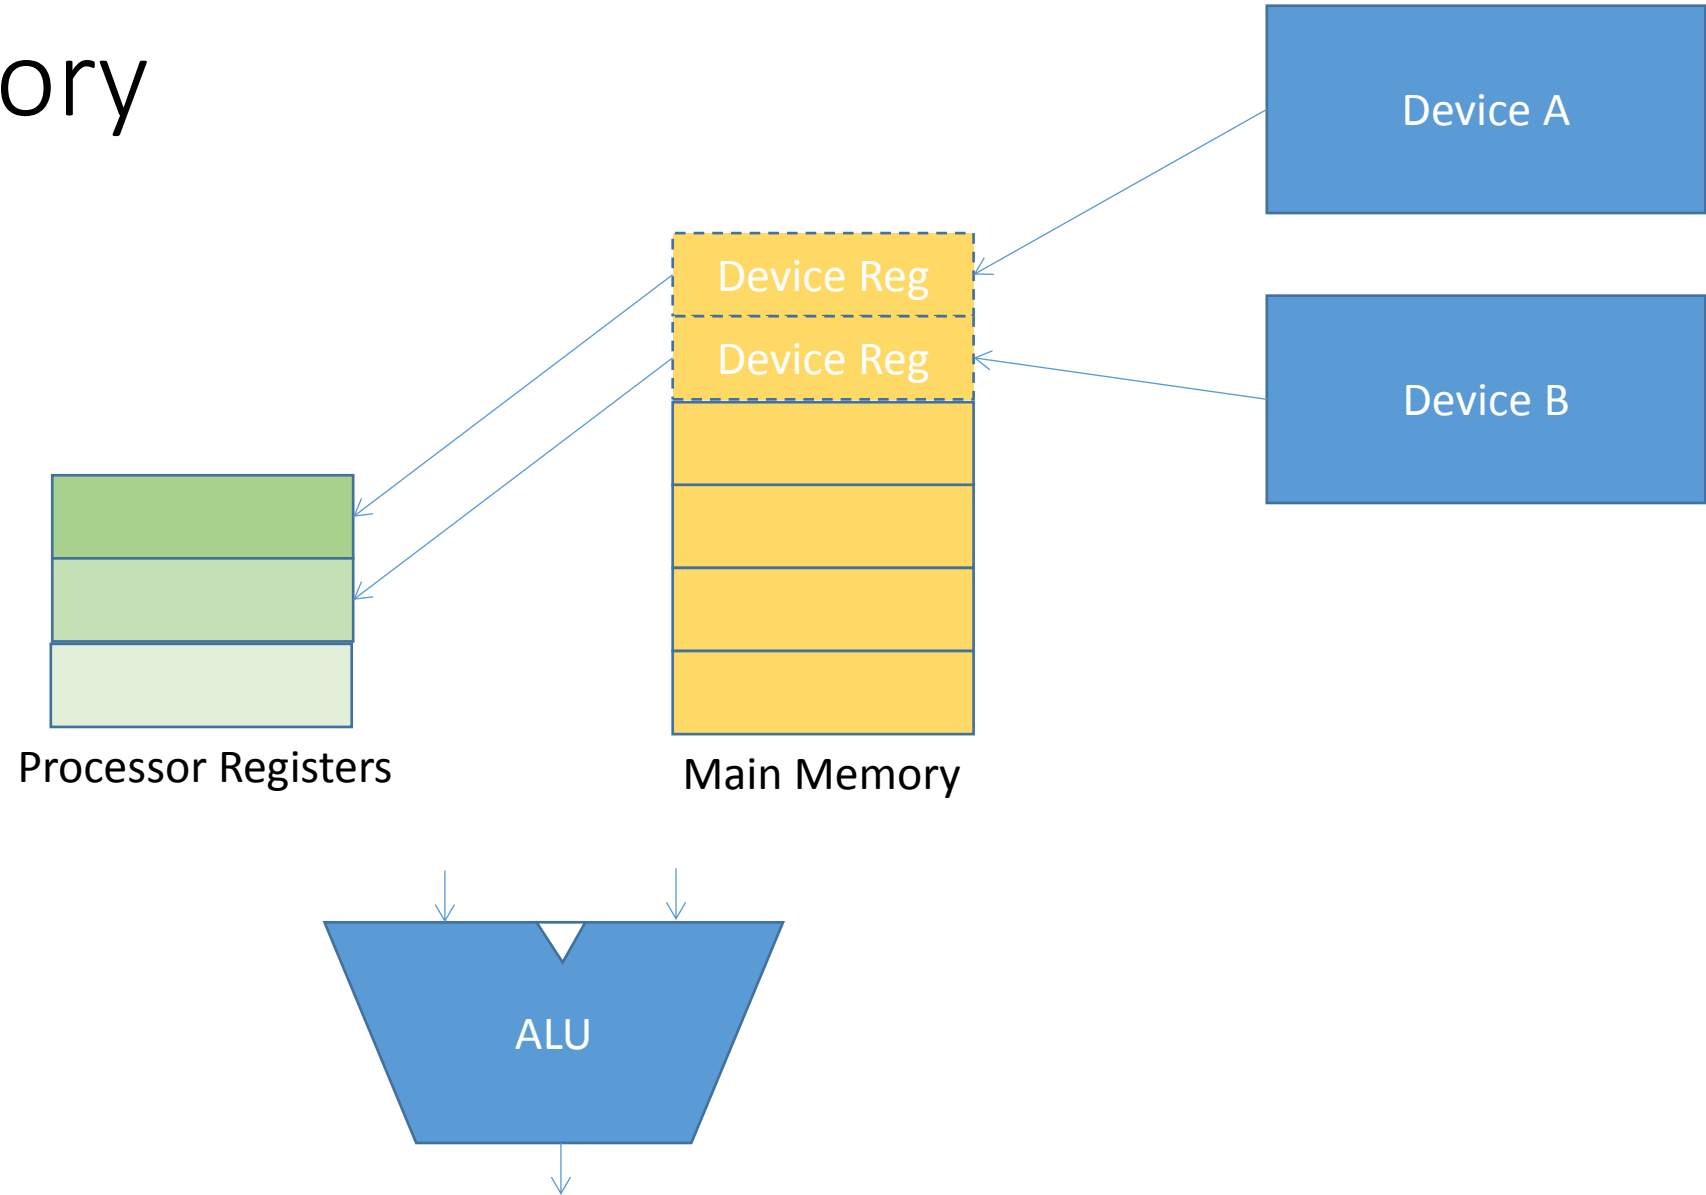
Conventional Memory



Conventional Memory



IO Memory



Memory Caching in IO Ports

- Volatile Keyword
- Disable hardware caching during IO access. Done by the kernel itself

Data Reordering

- Consider the Following Code Snippet
 - *Write (PIN FUNCTION, CLOCK DEVICE)*
 - *Write (FREQUENCY SETTING, CLOCK DEVICE)*
 - *Write (PRESCALAR VALUE, CLOCK DEVICE)*
 - *Write (START CLOCK, CLOCK DEVICE)*
- Compiler Reorders the data for optimized performance
 - *Write (PIN FUNCTION, CLOCK DEVICE)*
 - *Write (PRESCALAR VALUE, CLOCK DEVICE)*
 - *Write (START CLOCK, CLOCK DEVICE)*
 - *Write (FREQUENCY SETTING, CLOCK DEVICE)*

Barriers

- The code will be modified like this
 - *Write (PIN FUNCTION, CLOCK DEVICE)*
 - *Write (FREQUENCY SETTING, CLOCK DEVICE)*
 - *Write (PRESCALAR VALUE, CLOCK DEVICE)*
 - *BARRIER();*
 - *Write (START CLOCK, CLOCK DEVICE)*

Barriers in Linux Kernel

- The linux kernel provides 4 macros to cover all possible ordering needs
 - `void barrier(void)`
 - The kernel makes it a point that the compiler optimizations are absent across the barrier. The memory values present in the CPU registers are immediately stored in the memory
 - `void rmb(void);`
 - `void read_barrier_depends(void);`
 - `void wmb(void);`
 - `void mb(void);`
 - These functions insert hardware memory barriers in the compiled instruction flow. They are supersets of the barrier macro.
 - An rmb guarantees that any reads appearing before the barrier are completed prior to the execution of any subsequent read.
 - wmb orders write and mb orders both read and write.
- Header - `#include <asm/system.h>`

Device Addressing


- IO Mapped IO or IO Ports

- Peripheral devices are mapped into IO space and are accessed through special instructions
- Additional control lines are required to segregate IO access
- Memory is better utilized


- Memory Mapped IO or IO Memory

- Peripheral devices are mapped into the system memory
- All access to this device space is performed using similar instructions that access memory
- Effective memory gets reduced

IO mapped IO (IOIO) or IO Ports



Memory Mapped IO (MMIO) or IO Memory



Communicating with IO Ports

- Drivers communicate with many devices through the I/O ports.
- Linux introduces a set of functions that may be used to gain exclusive access to an I/O region and communicate data transfers to and from these ports.
- Before embarking into communication with the I/O ports, the driver must gain access to the required ports by calling the function
 - `struct resource *request_region(unsigned long first, unsigned long n, const char *name);`
 - *The argument first is the first address requested*
 - *The argument long specifies the length of the addresses requested*
 - *The argument name identifies the requested region by this name in /proc/ioports*
- Header - `#include <linux/ioport.h>`

Communicating with IO Ports

- When the I/O ports have been used as per requirement, they should be returned to the kernel so that other drivers may avail their existence. The I/O ports are returned by the function
 - `void release_region(unsigned long start, unsigned long n);`
- Before using a I/O region, a check on the availability has to be conducted to be certain that the requested region will be available for our use. A special function allows this check.
 - `int check_region(unsigned long first, unsigned long n);`
 - *This function returns a negative error code if the region is not available. This is a deprecated function and may not always prove to be true since it does not run atomic with the request_region function.*

Communicating with IO Ports

- On successfully being allotted the region, the actual communication to the ports may be carried out using the kernel provided functions.
- These functions are specific to the port sizes on the I/O devices and provide interfaces for 8-bit, 16-bit and 32-bit ports
 - `unsigned inb (unsigned port);`
 - `unsigned inw (unsigned port);`
 - `unsigned inl (unsigned port);`
 - *Get data from the port specified as an argument*
 - `void outb (unsigned char byte, unsigned port);`
 - `void outw (unsigned short word, unsigned port);`
 - `void outl (unsigned long word, unsigned port);`
 - *Send data to the port specified in the argument*
- Header - `#include <asm/io.h>`

Communicating with IO Ports

- The kernel also supports string operations that render its services in allowing a string of 'n' bytes to be transferred between the driver and the I/O port
 - `void insb(unsigned port, void *addr, unsigned long count);`
 - Similarly..`insw, insl`
 - *Copies count bytes of information from the port to addr*
 - `void outsb(unsigned port, void *addr, unsigned long count);`
 - *Write count data pointed by addr to the port*
- The kernel provides mechanisms of synchronization between a high end processor and a relatively slower I/O by allowing a pause functionality in data transfer.
- This feature may be accessed by ending the function names previously discussed with an '_p', such as `inb_p`, `outb_p`....

IO Memory

- IO Memory regions must be allocated prior to use
 - `struct resource *request_mem_region(unsigned long start, unsigned long len, char *name);`
- IO Memory regions should be released when no longer in use
 - `void release_mem_region(unsigned long start, unsigned long len);`
- Accessing IO Memory regions should be preceded by a call to
 - `void *ioremap(unsigned long phys_addr, unsigned long size);`
 - `void *ioremap_nocache(unsigned long phys_addr, unsigned long size);`
- After accessing the IO Memory region, unmap it using
 - `void iounmap(void * addr);`

IO Memory

- To read from IO memory,
 - unsigned int ioread8(void *addr);
 - unsigned int ioread16(void *addr);
 - unsigned int ioread32(void *addr);
- Writing to the IO memory through
 - void iowrite8(u8 value, void *addr);
 - void iowrite16(u16 value, void *addr);
 - void iowrite32(u32 value, void *addr);
- Reading a series of values
 - void ioread8_rep(void *addr, void *buf, unsigned long count);
 - void iowrite8_rep(void *addr, const void *buf, unsigned long count);

Interrupts - Definition

- Means to notify the kernel that a device is requesting attention.
- Characteristics of Interrupts
 - Interrupts occur Asynchronously
 - Processor completes the current instruction being executed
 - Jumps to Interrupt Service Routine
 - Handle the Interrupt
 - Return from Interrupt
 - Global Interrupts are disabled. It may be enabled in the handler by the driver
 - Interrupts must be handled quickly because the longer interrupts take to execute, the longer interrupts remain disabled
 - No Sleep, or other delay functions should be called in an interrupt
 - Less important functions of the interrupt should be performed in a bottom half (tasklet)

Interrupts

- Components

- IRQ Line – Interrupt Number

- These are specific to a peripheral
 - It could be shared between number of devices
 - Limited number (32 in x86). You have to find out your IRQ number before you can use it. It is very specific to the hardware that you are using

- Interrupt Handler

- Each handler has to register itself with the kernel whenever an operation is to be performed on the device.
 - Registering a handler is a notification by the driver to the kernel, claiming authority for access to the device through a requested IRQ number.
 - On completion of access to the device, the handler may be unregistered and the irq number freed for allowing access to other applications.

Interrupts – Requesting an IRQ and binding a handler

- Registering an interrupt handler to an irq number and notifying the kernel is done by
 - `int request_irq (unsigned int irq,
irqreturn_t (*handler) (int, void *, struct pt_regs *),
unsigned long flags,
const char *dev_name,
void *dev_id);`
 - The value returned from request_irq to the is either 0 to indicate success or a negative error code, as usual.
 - It is not uncommon for the function to return -EBUSY to signal that another driver is already using the requested interrupt line.
- The handler is freed by calling
 - `void free_irq (unsigned int irq, void *dev_id);`

Interrupts – Where to request for IRQ

- There are two places where an IRQ can be requested.
 - During the driver initialization at `module_init` function
 - For this, it is advisable that the interrupt is requested as a shared IRQ
 - In the open method of the driver
 - If called in the open call, it allows the use of the interrupt line only when the application requests for it.
 - The disadvantage of this technique is that a per device open count has to be maintained to know when interrupts can be disabled.

Interrupts – FLAGS passed to request_irq

- **FLAGS**
 - **SA_INTERRUPT** (May be deprecated.. Just check)
 - This bit indicates a “fast” handler.
 - Fast interrupt handlers run with all interrupts disabled on the local processor
 - **SA_SHIRQ**
 - This bit indicates that the interrupt can be shared between devices.
 - The dev_id must be unique to each registered handler. A pointer to any per-device structure is sufficient. NULL cannot be passed to this field
 - The interrupt handler must be capable of detecting whether its device generated an interrupt. This requires both hardware support and associated logic in the handler
 - **SA_SAMPLE_RANDOM**
 - Adds to the kernel entropy pool to increase the randomness.

Interrupts – The Handler Function

- The Interrupt Handler

- `static irqreturn_t intr_handler (int irq, void *dev_id, struct pt_regs *regs);`
 - *The return value of an interrupt handler is the special type `irqreturn_t`.*
 - *An interrupt handler can return two special values, `IRQ_NONE` or `IRQ_HANDLED`.*
 - *`IRQ_NONE` is returned when the handler detects an interrupt for which its device was not the originator.*
 - *`IRQ_HANDLED` is returned if the interrupt handler was correctly invoked*

Interrupts – Disabling and Re-enabling

- Single Interrupts

- Sometimes the driver may need to disable interrupt delivery for a specific line. This is achieved by using one of the functions
 - `void disable_irq(int irq);`
 - `void enable_irq(int irq);`


- Disable all Interrupts

- `void local_irq_save(unsigned long flags);`
- `void local_irq_disable(void);`
 - shuts off the interrupts
- `void local_irq_restore(unsigned long flags);`
- `void local_irq_enable(void);`
 - Re-enables the interrupts

Doing it on the Beagle Bone Black

Beaglebone Black Pinout Diagram

| P9 | | | |
|-----------|---------------|----|-----------|
| Function | Physical Pins | | Function |
| DGND | 1 | 2 | DGND |
| VDD 3.3 V | 3 | 4 | VDD 3.3 V |
| VDD 5V | 5 | 6 | VDD 5V |
| SYS 5V | 7 | 8 | SYS 5V |
| PWR_BTN | 9 | 10 | SYS_RESET |
| UART4_RXD | 11 | 12 | GPIO_60 |
| UART4_TXD | 13 | 14 | EHRPWM1A |
| GPIO_48 | 15 | 16 | EHRPWM1B |
| SPI0_CS0 | 17 | 18 | SPI0_D1 |
| I2C2_SCL | 19 | 20 | I2C2_SDA |
| SPI0_DO | 21 | 22 | SPI0_SLCK |
| GPIO_49 | 23 | 24 | UART1_TXD |
| GPIO_117 | 25 | 26 | UART1_RXD |
| GPIO_115 | 27 | 28 | SP11_CS0 |
| SP11_DO | 29 | 30 | GPIO_112 |
| SP11_SCLK | 31 | 32 | VDD_ADC |
| AIN4 | 33 | 34 | GND_ADC |
| AIN6 | 35 | 36 | AIN5 |
| AIN2 | 37 | 38 | AIN3 |
| AIN0 | 39 | 40 | AIN1 |
| GPIO_20 | 41 | 42 | ECAPWMO |
| DGND | 43 | 44 | DGND |
| DGND | 45 | 46 | DGND |



| LEGEND | |
|------------------------|--|
| Power, Ground, Reset | |
| Digital Pins | |
| PWM Output | |
| 1.8 Volt Analog Inputs | |
| Shared I2C Bus | |
| Reconfigurable Digital | |

| P8 | | | |
|------------|---------------|----|-------------|
| Function | Physical Pins | | Function |
| DGND | 1 | 2 | DGND |
| MMC1_DAT6 | 3 | 4 | MMC1_DAT7 |
| MMC1_DAT2 | 5 | 6 | MMC1_DAT3 |
| GPIO_66 | 7 | 8 | GPIO_67 |
| GPIO_69 | 9 | 10 | GPIO_68 |
| GPIO_45 | 11 | 12 | GPIO_44 |
| EHRPWM2B | 13 | 14 | GPIO_26 |
| GPIO_47 | 15 | 16 | GPIO_46 |
| GPIO_27 | 17 | 18 | GPIO_65 |
| EHRPWM2A | 19 | 20 | MMC1_CMD |
| MMC1_CLK | 21 | 22 | MMC1_DAT5 |
| MMC1_DAT4 | 23 | 24 | MMC1_DAT1 |
| MMC1_DAT0 | 25 | 26 | GPIO_61 |
| LCD_VSYNC | 27 | 28 | LCD_PCLK |
| LCD_HSYNC | 29 | 30 | LCD_AC_BIAS |
| LCD_DATA14 | 31 | 32 | LCD_DATA15 |
| LCD_DATA13 | 33 | 34 | LCD_DATA11 |
| LCD_DATA12 | 35 | 36 | LCD_DATA10 |
| LCD_DATA8 | 37 | 38 | LCD_DATA9 |
| LCD_DATA6 | 39 | 40 | LCD_DATA7 |
| LCD_DATA4 | 41 | 42 | LCD_DATA5 |
| LCD_DATA2 | 43 | 44 | LCD_DATA3 |
| LCD_DATA0 | 45 | 46 | LCD_DATA1 |

Introducing GPIO on BBB

- Header - `#include <linux/gpio.h>`
- Functions
 - `static inline bool gpio_is_valid(int number)`
 - check validity of GPIO number (max on BBB is 127)
 - `static inline int gpio_request(unsigned gpio, const char *label)`
 - allocate the GPIO number, the label is for sysfs
 - `static inline void gpio_free(unsigned gpio)`
 - deallocate the GPIO line
 - `static inline int gpio_export(unsigned gpio, bool direction_may_change)`
 - make available via sysfs and decide if it can change from input to output and vice versa
 - `static inline void gpio_unexport(unsigned gpio)`
 - remove from sysfs




Introducing GPIO on BBB

- `static inline int gpio_direction_input(unsigned gpio)`
 - an input line (as usual, return of 0 is success)
- `static inline int gpio_direction_output(unsigned gpio, int value)`
 - value is the state
- `static inline int gpio_get_value(unsigned gpio)`
 - get the value of the GPIO line
- `static void gpio_set_value(unsigned gpio, value)`
 - Set the value of the GPIO line
- `static inline int gpio_set_debounce(unsigned gpio, unsigned debounce)`
 - set debounce time in ms (platform dependent)
- `static inline int gpio_sysfs_set_active_low(unsigned gpio, int value)`
 - set active low (invert operation states)
- `static inline int gpio_to_irq(unsigned gpio)`
 - associate with an IRQ

Case Study – 1: Variable Frequency LED Toggle

- Involves IOCTL, Kernel Timers and GPIO
- Toggle an LED connected to GPIO_49 of your Beaglebone in a specific periodicity
- Use Kernel Timers to maintain toggle periodicity
- Using IOCTL from the user space, change the frequency of toggling by sending out a command through the driver to the LED
- Demonstrate...

Required Circuit



Case Study – 2: Toggle an LED when Switch is Pressed

- Connect a tactile switch to GPIO_115 on your BBB
- The switch should be pulled up, through a 15K Ohm resistor
- Whenever the switch is pressed, an interrupt should be generated and in the handler, the LED should be toggled
- Register the corresponding interrupt handler and IRQ number through GPIO access

Required Circuit

